

Authorization and Policy for AI Agents

Identity established who is acting. Now the platform has to decide, for every single tool call, whether this agent may perform this action, on this resource, in this context, on whose authority. It also has to decide whether a human must still say yes.

Part 5 of the Production AI Engineering series · Full edition · 41 min read ·

POC run 2026-09-29-recorded

F1 MCP Tool Sprawl

F2 Layered Architecture

F3 Headless AI

T1 Agent Identity


T2 Authorization & Policy

T3 Human-in-the-Loop

— PRODUCTION AI ENGINEERING · PART 5

Authorization & Policy for AI Agents

One identity. One request. Three possible answers.



Production AI Engineering · **T2 · Authorization & Policy** · builds on F1 MCP Tool Sprawl, F2 Layered Architecture, F3 Headless AI and T1 Agent Identity

Identity establishes who is acting. Authorization determines whether that identity may perform this specific action, on this specific resource, in this specific context. Approval determines whether it may proceed without a human. Three questions, three answers, three different parts of the platform.

This is the fifth article in the series. Here's where the others left off:

- **MCP Tool Sprawl** was the problem: every agent owned its own tools, connections and credentials.
- **Layered Architecture** was the structure: reasoning surrounded by clear boundaries for state, models, tools, workflow and control.
- **Headless AI** was reuse: the same intelligence invoked by events, services, workflows and other agents, not only by chat.
- **Agent Identity** was the chain: invoker → agent → runtime → credential → action, carried and recorded at every hop.

Agent Identity closed on a question it deliberately left open: *should incident-intel be allowed to roll back payment-service in production, at 14:09, during a SEV-1, with a change freeze starting at 15:00? And who decides?*

This article answers it.

14:02–14:07. Same Agent. Five Asks. Five Answers.

It is the same incident. At 14:02 a Datadog monitor fires: **payment-service error rate 14%, against a 1% SLO**. Release `v4.18.0` shipped 13 minutes earlier. The headless runtime wakes the incident agent.

The last article did the hard work of establishing who is acting. The platform now knows:

principal	incident-agent-prod	(registered as agent.incident-intel@1.3.0)
invoked by	svc.monitoring-webhook	
acting for	sre-team	(delegation d1g-sre-incident-2026q4)
runtime	spiffe://prod/agent-runtime	

With that established, the agent starts asking for things. Over the next few minutes it wants to read the deployment, restart pods, roll back staging, roll back production, and, following a stale runbook line pasted into the alert's annotation, delete a namespace.

```

Datadog Event
  ↓
Incident Agent
  ↓
Identity established
  ↓
Policy Check
  ↓
Can read deployment?    YES
Can restart pods?       YES , one pod per call
Can rollback staging?    YES
Can rollback production? CONDITIONAL : a human must approve
Can delete namespace?   NO

```

Every one of those requests comes from **the same identity**, the same agent, the same run and the same delegation. The identity is identical all the way down the list, and the answers are not.



Figure 1. Same identity, five asks, five answers of four kinds. Identity can't produce this list. Only policy can.

The difference between those answers has nothing to do with *who* is asking. It comes from **what** is being asked, **against what**, **where**, **when**, and **on whose authority**. That is the domain of authorization, and it is the subject of this article.

Why Identity Is Not Enough

Identity is a precondition for authorization, not a substitute for it. Knowing that the caller is `incident-agent-prod` tells you precisely nothing about whether it should delete `payments-canary`.

Many automation failures are not failures of identity at all. The automation was exactly who it said it was, it held a valid credential, and it did exactly what that credential permitted. The failure was that the credential permitted far too much, in far too many contexts, and nothing asked a second question.

Identity tells us who is acting. Authorization tells us what that identity may do.

The previous article separated five concerns that tend to be mashed together. This article adds a sixth, the one the next article covers, and puts authorization in the middle where it belongs.

Concern	The question it answers	Produces	In our incident
Authentication	Can you prove the claim you are making, right now?	A verified claim	The webhook signature checks out. The runtime's SVID is valid
Identity	Who are you, stably, across calls and time?	A principal	<code>incident-agent-prod</code> (<code>agent.incident-intel@1.3.0</code>)
Delegation	Whose authority are you borrowing, and how much?	A bounded grant	sre-team lends restart and rollback, for its services, during incidents
Authorization	Is <i>this</i> action on <i>this</i> resource permitted, <i>here and now</i> ?	A decision per call	Staging rollback: ALLOW. Namespace deletion: DENY
Approval	Even if permitted, must a human accept this exact call first?	A consent tied to one call	The incident commander approves the production rollback
Human-in-the-loop	When and how does a human intervene, redirect or stop the agent?	An operating model	<i>The next article</i>

Authentication ≠ identity ≠ ... ≠ approval.

Authorization sits between knowing who is acting and deciding whether a human must step in.




Figure 2. Authorization sits between knowing who is acting and deciding whether a human must step in.

Six distinctions run through the rest of this article. It is worth stating them plainly once, because every anti-pattern later on is a violation of one of them:

1. **Identity ≠ Authorization.** Knowing who is acting says nothing about what they may do.
2. **Tool access ≠ permission for every tool action.** Reaching the Kubernetes server is not permission to do everything Kubernetes can do.
3. **Authorization ≠ Approval.** Permitted in principle is not the same as cleared to execute now.
4. **Static permissions ≠ contextual decisions.** The same action can be fine at 14:07 and forbidden at 15:04.
5. **Hard-coded checks ≠ policy-driven control.** An `if` statement inside one agent is not a control the organisation can inspect, test or change.
6. **Agent autonomy ≠ unrestricted action.** Autonomy is about who *chooses* the action. It says nothing about how far the action may *reach*.

Four more confusions sit close to these, and the article corrects each where it comes up: an **OAuth scope is not fine-grained authorization** (it can't see context), **delegation is not impersonation** (the agent acts *for* someone, never *as* them), **a sandbox is not authorization** (authorization decides whether an action *may* happen, a sandbox whether the process *can* do it; you need both), and **a hash chain is not immutable storage** (it makes tampering evident, not impossible).

The Production Question: What Is This Agent Allowed to Do?

Here's the question as most teams first phrase it:

Can the agent use Kubernetes?

And here's the question production actually needs answered:

Can this agent perform this Kubernetes operation, against this resource, right now, on behalf of this principal?

The second question has more moving parts, and every one of them matters. Written out as an explicit request, it looks like this:

```
Can
principal:    incident-agent-prod
acting_for:   sre-team
action:       kubernetes.rollbackDeployment
resource:     deployment/payment-service
environment:  production
context:
  incident:    INC-4471 (SEV-1, open)
  change_window: normal
  diagnosis_evidence_score: 0.94 (computed by the platform)
?
```

That's the shape of every authorization request in this article: **principal, acting-for, action, resource, environment, context**. Remove any one of them and the decision becomes a guess. The POC carries exactly these fields in one structured `Request` object. The principal and acting-for come from the execution identity of the last article, the action, resource and arguments describe the proposed call, and the context is looked up by the platform. The `diagnosis_evidence_score` is not something the agent says about itself: the platform computes it from signals it observed. More on that below.

Why agents make this harder

Traditional services usually have a more predictable call graph: engineers decide most of their integrations and actions when the software is built. An agent is different. It chooses tools at run time, chains actions into plans, interprets untrusted context, and may exercise authority delegated by someone else. The authorization boundary has to survive decisions that were never enumerated in application code. Five properties make the difference.

Property	Traditional service	AI agent	What authorization must do
Tool selection	Fixed in code, reviewed before deploy	Chosen by a model at run time, from whatever it can see	Decide per call, not per deployment
Execution shape	One request, one action	Multi-step plans; each step conditioned on the last	Evaluate each step; earlier permission is not later permission
Authority	Its own, usually	Often borrowed: from a human, a team, another agent	Know whose authority each call spends
Inputs	Structured, validated	Natural language, alerts, logs, tickets, all potentially hostile	Never let input text widen what is permitted
Blast radius	Bounded by the code	Bounded only by the credentials it can reach	Bound it by policy, not by hope

The fourth row deserves a moment. At 14:04 in our incident, the agent proposes deleting `payments-canary`, not because it's malicious, but because the alert annotation contains a stale runbook line: *"if canary is unhealthy, kubectl delete ns payments-canary."* The model did what models do: it followed plausible instructions in its context. That is untrusted instruction (a mild, accidental prompt injection) meeting what OWASP's Top 10 for LLM Applications calls **Excessive Agency** (LLM06): more functionality, permission or autonomy than the task needs. You don't fix it by writing a better prompt. You fix it with a boundary the prompt cannot move.

Agent autonomy decides which action to attempt. Authorization decides which attempts are allowed to land.

Tool Access Is Not Authorization

This is the distinction that most agent platforms get wrong first, so it gets its own section.

After MCP Tool Sprawl, the platform has a capability gateway and a Kubernetes MCP server behind it. The agent's tool list includes Kubernetes. The tempting conclusion:

agent can access Kubernetes MCP server ⇒ agent can use Kubernetes

That arrow is the bug. The Kubernetes server exposes reads, restarts, rollbacks, scaling, exec, and deletion. They range from harmless to catastrophic, against dozens of namespaces in several environments. **Access to the server is access to the menu, not permission to order everything on it.**

Authorization for agents needs four levels of granularity, and each one narrows the one above:

Level	Question	Example	Who usually gets this right
Tool	May the agent reach this system at all?	incident-agent-prod may call the Kubernetes server	Almost everyone (MCP allowlists)
Action	Which operations on it?	read, restart, rollback: yes. deleteNamespace, exec: no	Some teams
Resource	On which objects?	payment-service (owned by sre-team), not checkout-service	Few teams
Context	Under which conditions?	production only with approval; not below 0.80 diagnosis evidence; not in a freeze unless SEV-1	Very few teams

Put the four levels together and the "Kubernetes" permission dissolves into something much more precise:

```
kubernetes.readPods           → ALLOW
kubernetes.readDeployment      → ALLOW
kubernetes.restartPod         → ALLOW (production: ALLOW_WITH_CONSTRAINTS , one pod per
call)
kubernetes.rollbackDeployment.stg → ALLOW
kubernetes.rollbackDeployment.prod → ALLOW_WITH_APPROVAL
kubernetes.exec               → DENY (no role grants it)
kubernetes.deleteNamespace    → DENY
```

Tool access is not authorization.

An MCP allowlist decides the top band. Production needs all four.




Figure 3. An MCP allowlist decides the top band. Production needs all four.

In production, tool access is not authorization.

MCP's own authorization specification is sometimes mistaken for the answer here. It isn't, and it doesn't claim to be. It is an OAuth-based way for a client to obtain a token that an MCP server will accept. That settles *which client is calling with which scopes*: the tool level, with some coarse action scoping. It doesn't, and can't, know that payment-service is in the middle of a SEV-1, that the change freeze starts at 15:00, or that sre-team does not own checkout-service. Those facts live elsewhere, and the decision has to be made somewhere that can see them.

From RBAC to Contextual Authorization

There are three well-established ways to model authorization. Each answers a different sub-question, and an agent platform needs all three. What follows frames them in agent terms, not as an IAM lecture.

RBAC: what the role normally permits

Role-based access control assigns permissions to roles and roles to principals. For agents, the role is the **ceiling**: the most this kind of agent may ever do, before any context is considered.

```
[roles.incident-responder]
actions = [
  "kubernetes.readPods", "kubernetes.readDeployment",
  "logs.query", "dashboards.read",
  "kubernetes.restartPod", "kubernetes.rollbackDeployment",
]

[roles.release-observer]
actions = ["kubernetes.readDeployment", "dashboards.read"]
```

TOML

`kubernetes.deleteNamespace` appears in no agent role. With deny-by-default, that's the end of the conversation for namespace deletion, whatever the prompt says.

RBAC is necessary and easy to reason about. On its own it's far too coarse: it can't tell staging from production, SEV-1 from SEV-4, or payment-service from checkout-service without a combinatorial explosion of roles (`incident-responder-prod-payments-sev1` ...).

ReBAC: what relationship the actor has to the resource

Relationship-based access control, the model popularised by Google's Zanzibar and its open-source descendants such as OpenFGA and SpiceDB, decides based on relationships in a graph: *sre-team owns payment-service*. *incident-agent-prod acts for sre-team*.

For agents, ReBAC answers the question RBAC can't: **whose authority is this, and does that authority extend to this particular object?**

```
sre-team      —owns—> service:payment-service
checkout-team —owns—> service:checkout-service
sre-team      —delegates—> incident-agent-prod
                                   (restart, rollback · while incident open)
```

At 14:07 the agent notices checkout-service is also erroring and proposes rolling it back too. Its role permits rollbacks. The delegation from sre-team is valid. But sre-team doesn't own checkout-service, so **it has no authority over checkout-service to lend**. The request is denied. Not because the action is dangerous, but because the relationship isn't there.

ABAC: how context changes the decision

Attribute-based access control evaluates rules over attributes of the principal, action, resource and environment. For agents, ABAC is where the **operational context** lives: environment, severity, time, change windows, diagnosis evidence, data classification, tenant boundaries.

```
[[rule]]
id = "F3-change-freeze"
effect = "DENY"
reason = "Change freeze: only SEV-1 remediation may write to production"
[rule.when]
environment = "production"
kind = "write"
change_window = "freeze"
"severity.not_in" = ["SEV-1"]
```

TOML

ABAC is where agent authorization becomes genuinely contextual, and also where it becomes easy to make a mess. A good discipline, which the POC enforces: **context rules only restrict**. They can deny, require approval, or add constraints. They can never grant an action no role grants.


How the three compose

The three models aren't competitors. They are three filters, applied in order, each answering one question:

Model	Agent-shaped question	Answers for the rollback	If it says no
RBAC	Does this kind of agent ever do this?	<code>incident-responder</code> permits <code>rollbackDeployment</code>	DENY: <i>no role grants it</i>
ReBAC	Whose authority is being spent, and does it cover this object?	sre-team delegated rollback; sre-team owns payment-service	DENY: <i>not delegated / not owner</i>
ABAC	Given everything happening right now, how should the answer change?	production → approval; SEV-1 during freeze → still allowed; evidence 0.94 ≥ 0.80	DENY, or ALLOW_WITH_APPROVAL, or ALLOW_WITH_CONSTRAINTS

Three lenses, one decision.


RBAC sets the ceiling, ReBAC finds the authority, ABAC applies the moment.



Three models are not competitors. They are three filters, applied in order.

Figure 4. Three lenses, one decision. RBAC sets the ceiling, ReBAC finds the authority, ABAC applies the moment.

In production the three lenses are often **different engines**, because they answer different kinds of question. A relationship graph answers "does sre-team own payment-service, and did it delegate to this agent?". That is what Zanzibar-style systems such as OpenFGA and SpiceDB are built for; OpenFGA's own documentation now models agents and acting-on-behalf-of relationships. A policy evaluator answers "given the role and this context, what's the effect?". That is Cedar's or OPA's job. The POC keeps all three in one small file, but the shape is worth seeing:



The context that matters for agents

These are the attributes that changed decisions in our incident, and where each one comes from. The second column matters more than it looks.

Attribute	Example	Source (who asserts it)	Why it matters
Environment	production / staging	Resource catalog	Same action, different consequence
Severity	SEV-1	Incident system	Emergencies justify more, faster
Incident state	open / mitigating / resolved	Incident system	Delegated incident authority should end with the incident
Change window	normal / freeze	Change calendar	Freezes protect the business from well-meaning changes
Diagnosis evidence	0.62 → 0.94	Evidence evaluator (platform-observed signals)	A weak diagnosis should not touch production
Data classification	restricted logs	Resource catalog	Reading may be fine; reading raw card numbers is not
Tenant boundary	all tenants / <code>acme</code> only	Incident system	An incident for one customer is not access to all customers
Blast radius	one pod vs every pod	Request arguments + policy	Shape the action, not just permit it

Notice what's missing from the second column: **the agent**. Nothing in this table is something the agent says about itself.

That is deliberate, and the diagnosis evidence row is where it matters most. The obvious design is to ask the model how confident it is and gate production writes on the answer. Don't. Self-reported model confidence shifts with phrasing and calibration, and an agent whose own number can move a production write from DENY to "eligible for approval" is grading its own homework. So the POC computes a **diagnosis evidence score** from signals the platform observed itself, in systems of record, before the decision:

Signal	Weight	Observed how
Release correlation	0.35	Production changed within 30 minutes of the alert, and the call reverts exactly that change
Error signature	0.27	Logs the gateway returned name the suspect release
Staging validation	0.22	The same fix was applied in staging, through the gateway
Staging recovery	0.10	Staging health recovered afterwards, per monitoring

At 14:05 the platform had seen the first two: **0.62**. By 14:07, after the staging rollback and recovery, all four: **0.94**. The weights are a design choice you'd tune. What matters is who produces the number.

Never let the agent grade its own homework. Every attribute that can loosen a decision must be looked up or observed by the platform, not asserted by the caller.

Static permissions versus contextual decisions

To show what "contextual" buys you, the POC takes **one request**, rolling back payment-service to v4.17.2, and evaluates it under eight different contexts. The agent, the role, the tool and the arguments are identical every time. Each row varies exactly one thing; everything else, including full diagnosis evidence, is held constant.

Context	Decision	Rule	Reason
staging	ALLOW	–	incident-responder permits kubernetes.rollbackDeployment
production, SEV-1, evidence 0.94	ALLOW_WITH_APPROVAL	A1-production-rollback	Production rollback requires SRE approval
production, evidence 0.62 (before staging validation)	DENY	F2-insufficient- evidence-high-risk-write	High-risk production writes need a diagnosis evidence score of at least 0.80
production, SEV-3, during change freeze	DENY	F3-change-freeze	Change freeze: only SEV-1 remediation may write to production
production, SEV-1, during change freeze	ALLOW_WITH_APPROVAL	A1-production-rollback	Production rollback requires SRE approval
production, incident resolved	DENY	rebac.delegation- inactive	Delegation dlg-sre-incident-2026q4 applies only while an incident on payment-service is open
production, no acting- for principal	DENY	rebac.no-principal	Writes need an acting-for principal; the agent has no write authority of its own
production, checkout- service (not owned)	DENY	rebac.not-owner	sre-team does not own checkout-service, so it cannot lend authority over it

Generated by the POC (`runs/2026-09-29-recorded/sweep.md`). The eight contexts are rows of one table in `scenario.toml` , each with its expected decision and rule; the same table produces `sweep.json` and drives the test for AUTHZ-INV-10, so the published table and the tested table can't drift apart.

One request, eight contexts: **1 ALLOW, 2 ALLOW_WITH_APPROVAL, 5 DENY**. A static permission model has exactly one answer for all eight rows. Whichever answer you pick, it's wrong for most of them.




Figure 5. The context sweep from the POC. A static permission can't be right for all eight rows.

Delegated Authority: Acting For Someone

Every write an agent performs spends somebody's authority. **Every write needs an attributable authority chain:** a record of whose authority is being spent, and proof that it covers this call.

This article models three sources of that authority:

- 1. The agent's own authority.** The agent is a principal with its own grants. This is fine for reading telemetry. It's dangerous for writing to production, because nobody accountable decided that this agent should change production today.
- 2. A human's authority.** Maya, the on-call SRE, invokes the agent and lends it part of her authority. This is delegation in the RFC 8693 sense: the agent keeps its own identity and acts *for* Maya, never *as* Maya. Her credentials are never forwarded into the agent runtime; the record of what she lent is.

Acting for someone is an intersection.

Every extra party can only take authority away.




Figure 6. Delegated authority is an intersection, never a union. Every extra party can only take authority away.

Two failure modes live at the edges of that intersection.

Borrowing more than the lender has. If sre-team doesn't own checkout-service, it can't lend authority over checkout-service, however broadly the delegation is worded. That is exactly what happened at 14:07 in the POC: the agent's role and delegation both covered rollbacks, and the decision was still DENY, because the relationship wasn't there.

The confused deputy. A low-privilege caller, whether another agent, a ticket or an injected instruction, asks a high-privilege agent to act. If the agent acts on its *own* broad authority, it has become a deputy for a caller that could never have done this itself. The fix is the same as in the last article: carry the acting-for principal through every hop, and authorize against the *principal whose authority is being spent*, not against the agent's own identity.

Delegated authority must be explicit, bounded, time-limited and auditable. If you can't point to the record that lent the authority, the agent didn't have it.

Policy Decisions in Practice

Now the architecture. Where do these decisions actually get made, and who enforces them?

Why the check can't live inside the agent

The tempting first implementation is a few `if` statements in the agent's tool wrapper:

```
if env == "production" and action == "rollback" and not approved:
    raise PermissionError("needs approval")
```

PYTHON

It works for one agent. Then there are six agents, written by four teams, each with its own subtly different version of that check. A change-freeze rule has to be edited in six repositories. Nobody can answer "what may agents do in production?" without reading all of them. And the check sits inside the very component whose behaviour you can't predict: a prompt-injected agent can simply call the tool by another path.

Policy has to be:

- **Shared:** one definition of "production rollback needs approval", used by every agent.
- **External to the agent:** the agent proposes, something else decides.
- **Inspectable:** someone other than the agent's author can read it and understand it.
- **Testable:** a change to policy runs a test suite before it ships, like any other code.
- **Changeable without redeploying agents:** a new freeze window is a policy change, not a release.
- **Enforced at one choke point:** the place every tool call already passes through.

Hard-coded checks are opinions scattered across codebases. Policy is a control the organisation owns.

PEP, PDP, PIP: three roles, briefly

The access-control world has a vocabulary for this (from XACML, and reused in NIST's Zero Trust architecture). It's worth borrowing three terms, and only three:

Role	What it does	In the agent platform
PEP: Policy Enforcement Point	Intercepts the call, asks for a decision, enforces it	The capability gateway from F1, the one every tool call already passes through
PDP: Policy Decision Point	Evaluates policy against the request and returns a decision	A policy evaluator (Cedar, OPA) for roles and context, often alongside a relationship graph (OpenFGA, SpiceDB) for ownership and delegation
PIP: Policy Information Point	Supplies the attributes the PDP needs	The systems of record : directory, relationship store, incident system, change calendar, resource catalog, and the evidence evaluator

Enforce · Decide · Inform.

The gateway enforces, the policy engine decides, systems of record supply the facts. The agent does none of the three.




Figure 7. The gateway enforces, the policy engine decides, and systems of record supply the facts. The agent does none of the three.

The division of labour is strict. **The PDP decides and never executes.** It doesn't call Kubernetes, doesn't ask a human, doesn't trust the agent to do anything. **The PEP enforces and never decides.** It doesn't interpret policy, it applies decisions. **The agent proposes and does neither.**

One property makes the whole arrangement real: **the gateway must be non-bypassable.** Otherwise every control in this article is advisory, and a security reviewer's first question, "what stops the model calling the Kubernetes API directly?", has no answer. It takes four things together:

Layer	Property	In the POC
Agent runtime	Holds no standing privileged credentials	Modeled: the agent only ever talks to the gateway
Credential broker	Issues a credential only after a permitting decision, scoped to one call	Modeled and tested: single-use, 60-second, signed, bound to the action, resource, environment and arguments

Layer	Property	In the POC
Tools	Refuse any call without a valid credential for exactly that call	Modeled and tested: direct, forged, mis-scoped, expired and reused credentials are all refused
Network	No route from the runtime to the protected APIs except through the gateway	Not modeled. It's network policy in production, and the POC doesn't claim it

XACML names a fourth role too, the PAP (policy administration point): where policy is written, reviewed and versioned. In this architecture that is the policy repository and its test suite.

Four decisions, not two

A yes/no answer is not expressive enough for production agents. **Our capability gateway exposes four operational outcomes.** The underlying engine may well return something simpler: Cedar returns allow or deny plus diagnostics, and XACML returns a decision plus *obligations* the PEP must fulfil. The gateway turns those into four enforceable execution states, from least to most restrictive. Call it the gateway's decision envelope:

Decision	Meaning	What the PEP does	Example
	Permitted as requested	Execute	Read the deployment
	Permitted, but only in a narrower shape	Reshape the call to fit, then execute	Restart one pod, not three; return logs redacted
	Permitted in principle; a named human role must accept this exact call first	Park the call, open an approval, stop	Roll back production
	Not permitted in this context	Block, explain, log	Delete a namespace

When several rules apply, **our gateway uses an explicit fail-closed combining rule: DENY → APPROVAL → CONSTRAINED → ALLOW.** That's our design choice, not a universal standard: XACML makes the combining algorithm configurable, and Cedar fixes it as *forbid overrides permit*. Constraints from several rules merge.

Unknown context fails closed, everywhere it can occur:

What is missing or unknown	What happens	Where it is tested
Principal, action, resource or environment not in the catalog	DENY (<code>default.unknown-*</code>)	AUTHZ-INV-01
Incident severity for a write	DENY (<code>default.unknown-incident-context</code>)	AUTHZ-INV-09

What is missing or unknown	What happens	Where it is tested
Ownership, delegation, or a delegation that expired, was revoked or is inactive	DENY (<code>rebac.*</code>)	AUTHZ-INV-04, 06, 09
Change calendar, or the evidence score	Any rule conditioned on it applies, so the restrictive outcome wins	AUTHZ-INV-09
The policy engine itself: unreachable, or an answer that isn't one of the four outcomes	The gateway denies (<code>POLICY_UNAVAILABLE</code>) and the tool is never called	L3 gateway tests

Constraints are underrated. At 14:05 the agent asks to restart three crash-looping pods in production. The decision is `ALLOW_WITH_CONSTRAINTS` with `maxPods: 1`. The gateway, not the agent, reshapes the call, restarts one pod, and records `maxPods=1 (requested 3)`. The agent gets a truthful response and re-plans. Constraints let policy say "yes, but smaller", which is often exactly the right answer for an agent in an incident.

What a decision looks like

Every decision has two audiences, and they get different views.

The agent gets a small, actionable envelope: the decision, a stable code, a plain reason, the next step, and any constraints it must live within. Here is exactly what the agent received for the production rollback:

```
{
  "decision": "ALLOW_WITH_APPROVAL",
  "code": "APPROVAL_REQUIRED",
  "reason": "Production rollback requires SRE approval",
  "next": "Wait for the approval decision",
  "constraints": {
    "approverRole": "IncidentCommander",
    "expiresIn": "10m"
  }
}
```

JSON

The audit log gets everything: which grants made the action possible, which rules matched, every attribute and evidence signal, the relationship path, and the policy version. We'll see that record in the audit section.

The split matters most for denials. A good denial is an actionable prompt, **but only with information the caller is allowed to learn**. When the checkout rollback was refused, the agent was told

`NO_DELEGATED_AUTHORITY`: "No delegated authority covers this action on this resource. Escalate to the owning team." It was *not* told which team owns checkout-service, or which rule fired, or what would

have made it pass. Telling an untrusted caller "denied by rule X, which would allow this if you used account Y" is a map for getting around the policy. A bare `403 Forbidden` teaches it nothing. The envelope teaches it the next step, and nothing more.

The Authorization POC: Watch It Work Under Pressure

Everything so far has been the mental model: identity, the four levels, the three lenses, delegated authority, the decision envelope. Now watch all of it operate in one incident, under time pressure, with the agent proposing the actions. (In the POC the agent is scripted rather than model-driven, so the run is deterministic; the policy boundary doesn't care who wrote the proposal.)

The POC exists to prove one thing: **an incident agent that goes through policy before every tool invocation behaves safely in exactly the situations where identity alone would not.** It extends the Headless AI and Agent Identity POCs, with the same incident and the same identifiers. It adds a real decision point, enforcement point, evidence evaluator, credential broker, approval gate and hash-chained decision log. It's standard-library Python: about 1,700 lines of POC and 500 lines of tests, with no network and no external Kubernetes. The enterprise systems are simulated, and the agent is scripted rather than model-driven, because the subject is the policy boundary, not the reasoning. Every decision, identifier, timestamp and measured value below comes from the recorded run (`authz_poc/runs/2026-09-29-recorded`).

The authorization POC.

One event, one agent, one enforcement point, one decision engine, one audit log.




Figure 8. The POC: one event, one agent, one enforcement point, one decision engine, one audit log.

Component	File	Responsibility
Alert source	scenario.toml	The Datadog event, including the stale runbook annotation
Incident agent	scenario.toml (scripted)	Proposes ten tool calls over seven minutes. Reports nothing about itself
Identity context	principals.toml, relationships.toml	Agent principal, roles, acting-for delegation
Policy engine (PDP)	authz/pdp.py + policy.toml	RBAC → ReBAC → ABAC, fail-closed combining
Attribute sources (PIP)	authz/pip.py + catalog.toml	Roles, ownership, severity, incident state, change window, classification, risk
Evidence evaluator	authz/evidence.py	Diagnosis evidence from platform-observed signals
Tool gateway (PEP)	authz/pep.py	Enforces decisions, applies constraints, binds approvals to exact calls, re-checks before executing, denies when the PDP fails
Credential broker	authz/broker.py	Issues a single-use credential for one call, only after a permitting decision

Component	File	Responsibility
Execution tools	<code>authz/tools.py</code>	Simulated Kubernetes, logs, monitoring. They know nothing about policy, and refuse any call without a broker credential for exactly that call
Approval gate	<code>authz/approvals.py</code>	Opens only for ALLOW_WITH_APPROVAL. One call, one approver role, one deadline, one use
Audit log	<code>authz/audit.py</code>	Hash-chained decision log
Invariants	<code>authz/invariants.py</code>	The 14 named production authorization invariants, <code>AUTHZ-INV-01</code> to <code>AUTHZ-INV-14</code>
Tests	<code>tests/test_l1_policy.py</code> ... <code>test_l4_replay.py</code>	Four layers: policy units, invariants, gateway integration, recorded replay
Verification	<code>authz/verify.py</code>	One machine-readable summary: invariants, test layers, replay, audit chain, determinism

The run, replayed



Figure 9. Seven minutes, seven beats. The identity never changes. Everything else does.

In seven minutes the agent proposes ten tool calls. Reads go through. The log query comes back redacted. A stale runbook line in the alert pushes the agent to delete a namespace, and it is stopped cold. A first production rollback, on thin evidence, is refused with a pointer to validate in staging. The agent restarts one pod (not the three it asked for), validates the fix in staging, reaches for checkout-service too and is refused, and then asks again for the production rollback. This time there's enough evidence to make it *approvable*. A human approves it. Policy checks it again. It lands at 14:09, the same minute as in the Headless and Identity runs.

Three moments worth zooming into

T2 · THREE CASES FROM THE RECORDED RUN

Zoom in on three decisions.



all values from authz_poc/runs/2026-09-29-recorded/decisions.jsonl and config/relationships.toml

Figure 10. Three decisions from the run, each proving a different claim.

Case A: a prompt cannot grant authority. The runbook line in the alert annotation read "if canary is unhealthy, kubectl delete ns payments-canary", and the agent proposed exactly that. It was denied three times over. No agent role grants namespace deletion, sre-team never delegated it, and an explicit forbid rule names it. Any one of the three would have been enough. That's defence in depth expressed as policy: widen a role by mistake next quarter and the forbid rule still holds, and an invariant test fails before the change ships. The agent was told `ACTION_NOT_PERMITTED` and "Escalate to a human operator". It wasn't told which of the three layers fired.

Case B: same call, different evidence. The production rollbacks at 14:05:10 and 14:07:30 carry the **same request fingerprint**, `8fc070d3f4fec204`. It is the identical call: same agent, acting-for, action, resource, environment and arguments. At 14:05 the platform had observed two of four evidence signals (0.62), so the decision was **DENY**, with `PROD_WRITE_EVIDENCE_INSUFFICIENT` and the next step "Validate the fix in staging first". The agent did exactly that. At 14:07 the platform had observed all four (0.94), and the same call became **ALLOW_WITH_APPROVAL**. Identity never changed. The agent's opinion of itself was never consulted. The evidence changed.

Case C: authority can't be borrowed from someone who doesn't have it. The agent's role permits rollbacks, and sre-team's delegation covers rollbacks. It still couldn't roll back checkout-service, because the authority being spent belongs to sre-team, and sre-team doesn't own checkout-service.

(The evidence rule also objected: nothing tied checkout's errors to a release.) This is the confused-deputy guard. The agent was told `NO_DELEGATED_AUTHORITY` and "Escalate to the owning team", without learning who that team is.

The approval, and the re-check

The production rollback didn't execute at 14:07:30. It was parked, and an approval was opened for that exact call. Here are the last four lines of the run's `transcript.txt`, unedited:

```
14:07:30 kubernetes.rollbackDeployment    production -> pending_approval [ ALLOW_WITH_APPROVAL ]
evidence 0.94
14:07:45 approval apr-819ba4aa10 by incident-agent-prod -> The requester cannot approve its own
action
14:08:50 approval apr-819ba4aa10 by ic.dev -> approved
14:09:02 agent re-submits the approved call: binding ok, policy re-check -> executed
{'rolled_back_to': 'v4.17.2'}
```

Three details matter.

- **The agent tried to approve its own request and was refused.** Who may approve is itself an authorization question, and "the requester" is never the answer.
- **The approval is bound to the exact call, not to an intention.** It isn't attached to the English phrase "roll back payment-service". It covers fingerprint `8fc070d3f4fec204` (this principal, acting-for, action, resource, environment, incident and arguments), under policy version `authz-2026-09-29.2`, for decision `dec-696fd603fc`, until 14:17:30, and once. When the agent re-submits, the gateway recomputes the fingerprint and checks every binding field. Change `to_version`, the resource or the policy version and the approval no longer applies; AUTHZ-INV-13 checks each of those.
- **Policy runs again at time of use.** Ninety seconds passed between request and execution. If a freeze had started, the incident had closed, ownership had changed or the delegation had expired in that window, the re-check would have denied the call. AUTHZ-INV-13 makes each of those changes between approval and execution and confirms the tool is never called. The approval authorized the call. It didn't freeze the world.

How a decision is made

For every proposed tool call, the gateway runs the same flow:

```

proposed call
↓
build request  principal · acting-for · action · resource
               environment · arguments · context (clock and incident id)
↓
PIP           roles, ownership, delegation, severity, incident state,
               change window, classification, risk, tenant scope,
               diagnosis evidence (high-risk writes); each with its source
↓
PDP  0. deny by default    unknown principal, action, resource,
                           environment · write without incident
                           severity → DENY
    1. RBAC                no role grants it → DENY
    2. ReBAC (writes)      no acting-for · not delegated · not owner
                           · delegation inactive → DENY
    3. ABAC guard rules    DENY · ALLOW_WITH_APPROVAL
                           · ALLOW_WITH_CONSTRAINTS
    → fail-closed combining: DENY > APPROVAL > CONSTRAINED > ALLOW
↓
PEP  PDP error / invalid   DENY (POLICY_UNAVAILABLE)
    DENY                  block; tell the agent code, reason, next
    ALLOW / CONSTRAINED   reshape to the constraints, get a one-call
                           credential from the broker, execute
    APPROVAL              open an approval bound to this call, stop
↓
audit  one hash-chained record per decision, approval and execution

```

The whole ABAC policy fits on one screen. That's deliberate: policy nobody can read is policy nobody can review. Each rule carries two audiences: `reason` for the audit record, and `code`, `message` and `next` for the agent.

```
version = "authz-2026-09-29.2"
```

```
[[rule]]
```

```
id = "F1-no-namespace-deletion"
```

```
effect = "DENY"
```

```
reason = "Namespace deletion is never delegated to agents"
```

```
code = "ACTION_NEVER_DELEGATED"
```

```
message = "This action is not available to agents"
```

```
next = "Escalate to a human operator"
```

```
[rule.when]
```

```
action = "kubernetes.deleteNamespace"
```

```
[[rule]]
```

```
id = "F2-insufficient-evidence-high-risk-write"
```

```
effect = "DENY"
```

```
reason = "High-risk production writes need a diagnosis evidence score of at least 0.80"
```

```
code = "PROD_WRITE_EVIDENCE_INSUFFICIENT"
```

```
message = "Not enough independent evidence for this production change"
```

```
next = "Validate the fix in staging first"
```

```
[rule.when]
```

```
environment = "production"
```

```
risk = "high"
```

```
"evidence_score.lt" = 0.80
```

```
[[rule]]
```

```
id = "F3-change-freeze"
```

```
effect = "DENY"
```

```
reason = "Change freeze: only SEV-1 remediation may write to production"
```

```
code = "CHANGE_FREEZE"
```

```
[rule.when]
```

```
environment = "production"
```

```
kind = "write"
```

```
change_window = "freeze"
```

```
"severity.not_in" = ["SEV-1"]
```

```
[[rule]]
```

```
id = "A1-production-rollback"
```

```
effect = "ALLOW_WITH_APPROVAL"
```

```
reason = "Production rollback requires SRE approval"
```

```
code = "APPROVAL_REQUIRED"
```

```
[rule.when]
```

```
action = "kubernetes.rollbackDeployment"
```

```
environment = "production"
```

```
[rule.constraints]
```

```
approverRole = "IncidentCommander"
```

```
expiresIn = "10m"
```

```
[[rule]]
```

```
id = "C1-production-restart"
```

```
effect = "ALLOW_WITH_CONSTRAINTS"
```

```
reason = "Production restarts are limited to one pod per call"
```

```
[rule.when]
```

```
action = "kubernetes.restartPod"
```

```
environment = "production"
```

```
[rule.constraints]
```

```
maxPods = 1
```

```
# ... plus F4-tenant-boundary and C2-restricted-logs (see authz_poc/config/policy.toml)
```

The format is incidental. The same intent in Cedar reads like this. There, `forbid` always overrides `permit` and everything is denied by default. The approval requirement travels as an annotation the PEP reads when that policy is among the ones that decided:

CEDAR

```
@id("rollback-grant")
// The PEP reads this annotation when this policy decides
@obligation("approval:IncidentCommander:10m")
permit (
  principal in Role::"incident-responder",
  action == Action::"rollbackDeployment",
  resource
) when {
  resource.owner == context.actingFor && context.delegationActive
};

@id("F2-insufficient-evidence-high-risk-write")
forbid (principal, action, resource)
when {
  action in Action::"highRiskWrite" && // an action group
  resource.environment == "production" &&
  context.evidenceScore.lessThan(decimal("0.80")) // computed by the platform
};
```

Cedar and OPA return allow or deny. Richer outcomes are modelled as **obligations**, the XACML term: requirements attached to the deciding policies that the enforcement point must fulfil. That is exactly the job of the gateway's four-state envelope. Pick the engine your organisation can operate. The concepts are what matter.

Stripped of logging, the PDP is short enough to review in one sitting:

```

def evaluate(self, req: Request) -> Decision:
    attrs = self.pip.attributes(req)          # facts from systems of record
    denials, grants = [], []
    is_write = attrs.get("kind") in ("write", "destructive")

    # 0. Deny by default: unknown principal, action, resource,
    #    environment; a write without known incident severity
    if not self.pip.principal_known(req.principal):
        builtin("default.unknown-principal")
    if "kind" not in attrs:
        builtin("default.unknown-action")
    if not self.pip.resource_known(req.resource):
        builtin("default.unknown-resource")
    if not self.pip.environment_known(req.environment):
        builtin("default.unknown-environment")
    if is_write and not self.pip.severity_known(attrs.get("severity")):
        builtin("default.unknown-incident-context")

    # 1. RBAC: the ceiling
    roles = [r for r in attrs["principal.roles"]
              if req.action in self.pip.role_actions(r)]
    if roles:
        grants.append(f"role:{roles[0]}")
    else:
        builtin("rbac.no-grant")

    # 2. ReBAC: whose authority does a write spend?
    if is_write:
        self._delegated_authority(req, attrs, now, grants, builtin)

    # 3. ABAC: context can only restrict
    effect, constraints = ALLOW, {}
    for rule in self.policy["rule"]:
        if not rule_matches(rule, attrs): # unknown attribute: fail closed
            continue
        if rule["effect"] == DENY:
            denials.append(rule)
        else:
            constraints.update(rule.get("constraints", {}))
            effect = max(effect, rule["effect"], key=STRICTNESS.get)

    # audit gets every reason; the agent gets code, reason, next
    return Decision(DENY, ...) if denials else Decision(effect, ..., constraints)

```

And the gateway's approval path, which is where binding and re-checking live:

```

def execute_approved(self, approval_id: str, req: Request) -> dict:
    apr = self.approvals.pending.get(approval_id)
    problem = (
        "unknown approval" if apr is None else
        "approval already used" if apr["status"] == "used" else
        "approval not granted" if apr["status"] != "approved" else
        "approval does not cover this call"
        if req.fingerprint() != apr["request_fingerprint"] else
        "approval is for another incident"
        if req.context.get("incident_id") != apr["incident_id"] else
        "policy changed since approval"
        if self.pdp.policy["version"] != apr["policy_version"] else
        "approval expired" if now >= apr["expires_at"] else None
    )
    if problem:
        return {"status": "denied", "code": "APPROVAL_NOT_APPLICABLE", "reason": problem}

    d = self._decide(req)          # time of check is not time of use
    if d.decision != APPROVAL:     # a freeze started, the incident closed ...
        return {"status": "denied", **d.public()}
    out = self._execute(req, d, approval=apr) # one-call credential, then the tool
    self.approvals.consume(approval_id, now) # single use
    return out

```

`_decide` is the PDP call wrapped so that an exception, or an answer that isn't one of the four outcomes, becomes a DENY. And `_execute` is the only place a credential is minted: the broker refuses a DENY, and refuses an approval-gated call without a granted approval.

The complete recorded run

Want the raw run? Here are all ten decisions, with the audit-view reason for each exactly as `timeline.json` records it, and the evidence score where the rule needed one:

#	Time	Agent proposes	Env	Evidence	Decision	Reason (audit view)
1	14:03	<code>readDeployment</code> payment-service	prod	—	ALLOW	incident-responder permits kubernetes.readDeployment
2	14:03	<code>readPods</code> payment- service	prod	—	ALLOW	incident-responder permits kubernetes.readPods
3	14:04	<code>logs.query</code> payment-service	prod	—	ALLOW_WITH_CONSTRAINTS	Restricted logs are returned redacted
4	14:04	<code>deleteNamespace</code> payments-canary	prod	—	DENY	No role held by incident-agent-prod grants kubernetes.deleteNamespace; sre-team has not delegated kubernetes.deleteNamespace to incident-agent-prod; Namespace deletion is never delegated to agents

#	Time	Agent proposes	Env	Evidence	Decision	Reason (audit view)
5	14:05	rollbackDeployment payment-service	prod	0.62	DENY	High-risk production writes need a diagnosis evidence score of at least 0.80
6	14:05	restartPod ×3 payment-service	prod	–	ALLOW_WITH_CONSTRAINTS	Production restarts are limited to one pod per call
7	14:06	restartPod payment-service	staging	–	ALLOW	incident-responder permits kubernetes.restartPod
8	14:06	rollbackDeployment payment-service	staging	0.62	ALLOW	incident-responder permits kubernetes.rollbackDeployment
9	14:07	rollbackDeployment checkout-service	prod	0	DENY	sre-team does not own checkout-service, so it cannot lend authority over it; High-risk production writes need a diagnosis evidence score of at least 0.80
10	14:07	rollbackDeployment payment-service	prod	0.94	ALLOW_WITH_APPROVAL	Production rollback requires SRE approval

Totals from facts.json : 10 calls proposed · 4 ALLOW · 2 ALLOW_WITH_CONSTRAINTS · 1 ALLOW_WITH_APPROVAL · 3 DENY · 7 executed (including the production rollback, after approval) · 3 blocked · 0 namespaces deleted.




Figure 11. The whole run on one time axis.

Ten calls. Seven minutes. Four kinds of answer.

Every tool call the agent proposed, with the decision and the reason the policy engine gave.

TIME	AGENT PROPOSES	ENV	DECISION	REASON
14:03	read deploy payment-service	prod	ALLOW	incident-responder permits kubernetes.readDeployment
14:03	read pods payment-service	prod	ALLOW	incident-responder permits kubernetes.readPods
14:04	query logs payment-service	prod	CONSTRAINED	Restricted logs are returned redacted
14:04	delete ns payments-canary	prod	DENY	No role held by incident-agent-prod grants kubernetes.deleteNamespace +2 more reasons
14:05	rollback payment-service · evidence 0.62	prod	DENY	High-risk production writes need a diagnosis evidence score of at least 0.80
14:05	restart pod payment-service	prod	CONSTRAINED	Production restarts are limited to one pod per call
14:06	restart pod payment-service	staging	ALLOW	incident-responder permits kubernetes.restartPod
14:06	rollback payment-service · evidence 0.62	staging	ALLOW	incident-responder permits kubernetes.rollbackDeployment
14:07	rollback checkout-service · evidence 0.00	prod	DENY	sre-team does not own checkout-service, so it cannot lend authority over it +1 more reason
14:07	rollback payment-service · evidence 0.94	prod	APPROVAL	Production rollback requires SRE approval
TOTALS	4 × ALLOW 2 × CONSTRAINED 1 × APPROVAL 3 × DENY executed 7 blocked 3			

from runs/2026-09-29-recorded

same call, 2.3 minutes apart

Identity never changed. The answers did.

Figure 12. The decision board: every call, every decision, every reason.

The receipt

This is the audit record of the policy decision for the production rollback, as the POC wrote it to `decisions.jsonl` (only the `attributes` and `attribute_sources` blocks are abridged). It records whose authority was spent and by which record (`delegation_chain`), the exact arguments, where each fact came from (`attribute_sources`), and both the audit reason and what the caller was told (`agent_view`):

```

{
  "seq": 18,
  "time": "2026-09-29T14:07:30Z",
  "kind": "policy.decision",
  "record": {
    "decision_id": "dec-696fd603fc",
    "recheck_of": null,
    "principal": "incident-agent-prod",
    "acting_for": "sre-team",
    "action": "kubernetes.rollbackDeployment",
    "resource": "deployment/payment-service",
    "environment": "production",
    "arguments": { "to_version": "v4.17.2" },
    "request_fingerprint": "8fc070d3f4fec204",
    "policy_version": "authz-2026-09-29.2",
    "grants": [
      "role:incident-responder",
      "delegation:dlg-sre-incident-2026q4",
      "relation:sre-team owns payment-service"
    ],
    "delegation_chain": [
      {
        "from": "sre-team",
        "to": "incident-agent-prod",
        "delegation": "dlg-sre-incident-2026q4",
        "granted_by": "sre.lead",
        "actions": ["kubernetes.restartPod", "kubernetes.rollbackDeployment"],
        "expires": "2026-12-31T23:59:59Z",
        "relation": "sre-team owns service:payment-service"
      }
    ],
    "matched_rules": ["A1-production-rollback"],
    "attributes": {
      "severity": "SEV-1",
      "incident_state": "open",
      "change_window": "normal",
      "risk": "high",
      "evidence_score": 0.94,
      "evidence_signals": ["release_correlation", "error_signature", "staging_validation",
"staging_recovery"],
      "...": "(abridged)"
    },
    "attribute_sources": {
      "severity": "incident system",
      "incident_state": "incident system",
      "change_window": "change calendar",
      "evidence_score": "evidence evaluator",
      "...": "(abridged)"
    },
    "decision": "ALLOW_WITH_APPROVAL",
    "reason": "Production rollback requires SRE approval",
    "constraints": { "approverRole": "IncidentCommander", "expiresIn": "10m" },
    "agent_view": {
      "decision": "ALLOW_WITH_APPROVAL",
      "code": "APPROVAL_REQUIRED",

```

```

    "reason": "Production rollback requires SRE approval",
    "next": "Wait for the approval decision",
    "constraints": { "approverRole": "IncidentCommander", "expiresIn": "10m" }
  },
  "prev": "3e31d14e4763929d",
  "hash": "16d33cf912830387"
}

```

And the execution record, five entries later, which links back to the re-checked decision, the approval and the single-use credential the broker issued for this one call:

```

{
  "seq": 23,
  "time": "2026-09-29T14:09:02Z",
  "kind": "tool.executed",
  "record": {
    "decision_id": "dec-0b5c5dde6e",
    "action": "kubernetes.rollbackDeployment",
    "resource": "deployment/payment-service",
    "environment": "production",
    "arguments": { "to_version": "v4.17.2" },
    "constraints_applied": [],
    "approval_id": "apr-819ba4aa10",
    "approved_by": "ic.dev",
    "credential": { "id": "cred-c97ffe5cf8", "expires_at": "2026-09-29T14:10:02Z", "single_use":
true },
    "output": { "rolled_back_to": "v4.17.2" }
  },
  "prev": "825f601473589408",
  "hash": "63b9c60232f8b1cc"
}

```

The run verifies the chain (`audit_chain_valid: true`), then flips one DENY to ALLOW in a copy and verifies again (`tamper_detected: true`). Be precise about what that buys. **A hash chain makes accidental or partial tampering evident. It does not stop someone who can rewrite the whole chain.** In production, ship the trail as it's written to an independently controlled, append-only audit store, or anchor it there, so that the people who run the agent platform can't quietly rewrite its history.

How it is tested: four layers, fourteen invariants

A security POC that only shows the happy path proves nothing, so the test suite is biased towards what the agent *cannot* do. It runs with one command, `python3 -m unittest discover -s tests -v`, needs no network and no cluster, freezes the clock to the scenario's, and is organised in four layers:

Layer	File	What it proves
L1 Policy unit tests	<code>tests/test_l1_policy.py</code>	The PDP on its own: a decision table (request → expected decision and deciding rule), the freeze case, the combining precedence, condition semantics, the evidence evaluator, the request fingerprint, the caller-facing view
L2 Authorization invariants	<code>tests/test_l2_invariants.py</code>	The 14 named production invariants below, each made of named cases that each build a fresh world
L3 Gateway integration	<code>tests/test_l3_gateway.py</code>	The real enforcement path, with every simulated tool counting its effects: denied calls never reach a tool, constraints are applied by the gateway, approval-gated calls wait, stale authority is re-checked, every execution leaves a receipt, tools refuse calls without a broker credential, and a failing or nonsensical policy engine means DENY
L4 Recorded replay	<code>tests/test_l4_replay.py</code>	The incident re-run into a fresh directory and compared byte for byte with the recorded run, and every published fact re-derived from the generated artifacts rather than restated in the test

The 14 invariants are the chapter's contract. Their number is fixed; the cases under each one can grow.

Id	Invariant	Representative cases
AUTHZ-INV-01	Deny by default	unknown principal, action, resource or environment → DENY
AUTHZ-INV-02	Role ceiling cannot be exceeded	<code>deleteNamespace</code> DENY in every context, with every context rule removed, and with the delegation widened to include it
AUTHZ-INV-03	Tool access is not action permission	same tool, same identity: read executes, delete and exec are denied; a read credential can't be spent on a delete; calling the tool directly changes nothing
AUTHZ-INV-04	Resource authority is required	rollback checkout-service: role ✓, delegation ✓, ownership ✗ → DENY
AUTHZ-INV-05	Delegation is an intersection, never a union	a delegation wider than the role, a role wider than the delegation, both without the relationship: each capped by the narrowest
AUTHZ-INV-06	Delegation scope and lifetime are enforced	wrong action, wrong resource, wrong or missing acting-for, expired, inactive incident, revoked
AUTHZ-INV-07	Context may restrict, never manufacture authority	injected ALLOW and constraint rules cannot grant <code>deleteNamespace</code> or checkout-service; context rules add no grants
AUTHZ-INV-08	Agent assertions never loosen authorization	confidence 0.99, "urgency: critical", "trust me", a claimed evidence score of 1.0 or a claimed approver change nothing

Id	Invariant	Representative cases
AUTHZ-INV-09	Missing or unknown security context fails closed	unknown or missing severity, missing environment, missing ownership, calendar unavailable, no evidence source, unknown incident
AUTHZ-INV-10	Same identity may produce different contextual decisions	the context sweep, read from the same <code>scenario.toml</code> table that produces <code>sweep.json</code>
AUTHZ-INV-11	Production rollback cannot silently auto-execute	qualifying production rollback → ALLOW_WITH_APPROVAL, parked, tool untouched; staging → ALLOW
AUTHZ-INV-12	Approval cannot widen DENY	no denied call opens an approval; the gate refuses to open for a DENY; an approved call whose re-check is DENY doesn't run; self-approval, wrong role, another agent, an expired window and reuse are all refused
AUTHZ-INV-13	Authorization is rechecked at time of use	between approval and execution: arguments, resource or policy version changed; delegation expired; incident closed; freeze started; ownership changed
AUTHZ-INV-14	Every decision is reconstructable and replayable	every decision record answers who, for whom, what, against what, context, policy version, rules and outcome; every execution links to its decision, credential and approver; editing any record breaks the chain; replay reproduces every decision

Three of them are deliberate mutation tests, the cheapest way to show the architecture is layered rather than resting on one rule. AUTHZ-INV-02 deletes every context rule and `deleteNamespace` is still denied, because no role grants it. AUTHZ-INV-07 injects permissive rules and nothing new becomes possible. AUTHZ-INV-05 widens the delegation and the role still caps it.

`python3 -m authz.verify` checks the recorded run end to end. It evaluates the invariants against today's code and policy, runs the whole suite twice to confirm identical results, replays the scenario into a fresh directory and compares every output byte for byte, and re-hashes the audit chain from the file alone. This is what it printed for the published run:

```
{
  "run": "2026-09-29-recorded",
  "policy_version": "authz-2026-09-29.2",
  "invariants": { "total": 14, "passed": 14, "failed": 0 },
  "tests": { "L1": { "passed": 24, "failed": 0 }, "L2": { "passed": 90, "failed": 0 },
    "L3": { "passed": 19, "failed": 0 }, "L4": { "passed": 20, "failed": 0 } },
  "integration": "PASS",
  "replay": "PASS",
  "audit_chain": "PASS",
  "deterministic": true,
  "ok": true
}
```

JSON

The per-layer numbers count individual cases and grow as cases are added; the stable facts are the ones above them and below: 14 of 14 invariants, every layer green, a byte-identical replay, a verified chain.



Figure 13. Every decision, identifier, timestamp and measured value in the POC figures comes from the recorded run. The results strip is read from `tests.json` and `verification.json`.

Every result is browsable. `results/lab-console.html` is a generated Lab Console for the run. Every check (each proposed call, approval attempt, sweep context, test case and invariant) opens with its input, its output, both audiences of the decision, the audit records it left and the files it came from. It is built from the recorded files only (`python3 tools/build_lab_console.py`) and refuses to build if they disagree. And `docs/TRACEABILITY.md` maps each claim in this article to the policy rule or code that implements it, the invariant that tests it, the artifact that records it and the figure that shows it.

What this proves, and what it doesn't

Within the modeled architecture, the run and the tests show:

- **Identity alone could not have produced these outcomes.** Ten calls from one identity got four different kinds of answer.

- **Tool access is not authorization.** The agent could reach Kubernetes throughout, yet three Kubernetes calls were denied, and two others, a pod restart and a log query, were reshaped by the gateway.
- **Context changes decisions, and evidence, not self-assessment, is what changed.** The same call, with the same fingerprint, was denied at 0.62 and routed to approval at 0.94, because the platform observed the staging validation. Nothing the agent asserts about itself can loosen a decision.
- **Delegated authority is enforceable.** Authority sre-team didn't hold (checkout-service) couldn't be spent, and authority it did hold ends with the incident, the delegation's expiry or its revocation.
- **Denied calls never execute, and the tools can't be reached around the gateway.** A denied call never reaches a tool; a tool refuses any call without a single-use credential minted for exactly that call after a permitting decision; and a failing policy engine means DENY.
- **Authorization and approval are separate steps.** Policy said the rollback was valid in principle. A separate gate decided when it could proceed, refused the agent's own approval, bound the approval to one exact call, and policy checked again at time of use.
- **Every decision is reconstructable and replayable,** from a log whose tampering is evident.

What it doesn't prove, deliberately:

- **The network half of non-bypassability.** The POC models credentials and tool-side checks. That the runtime has no network route to the real APIs is a deployment property it can't demonstrate.
- **A real model, a real IdP or real systems of record.** The agent is scripted, the broker signs with a fixed key, and the incident system, calendar, catalog, Kubernetes and logs are simulated.
- **That these evidence weights are right for your systems,** or that this evaluator scales, or behaves under concurrency.
- **A protected audit store.** The hash chain lives in process; production must ship or anchor it elsewhere.
- **Human-in-the-loop.** Escalation, reminders, quorum, timeouts and intervention are the next article's.

A production platform would use a hardened engine (Cedar or OPA, plus OpenFGA or SpiceDB for relationships), real systems of record, a real token broker and a protected audit store. The boundary, and what crosses it, stays the same.

Where Approval Fits

This is the distinction the whole article has been building towards, so it's worth stating as sharply as possible.


Authorization determines whether an action is permitted. Approval determines whether it may proceed without a human.

They are different questions, answered at different times, by different parties, with different inputs.

	Authorization	Approval
Question	Is this action valid for this principal, on this resource, in this context?	Should this specific, valid action happen now?
Answered by	A policy engine, automatically, in milliseconds	A named human in an accountable role
Inputs	Roles, relationships, attributes, policy	The proposal, the evidence, the human's judgement
Scope	A class of requests, defined in advance	One exact call, with its arguments
Can it widen what's allowed?	It defines what's allowed	Never. Approval can only narrow
Output	ALLOW · CONSTRAINED · APPROVAL · DENY	Approved / rejected / expired

Two gates, not one.

Policy decides what is valid. A human decides whether a valid, high-risk action proceeds.



Authorization says the action is valid. Approval says it may proceed.

Figure 14. Two gates, not one. Policy decides what's valid. A human decides whether a valid, high-risk action proceeds.

Why keep them apart? Because merging them fails in both directions.

If approval can widen authorization, a human clicking "approve" becomes a way around policy. The agent asks for something no role permits, someone tired at 03:00 approves it, and the platform executes it. In the POC, approval is only ever reachable from an `ALLOW_WITH_APPROVAL` decision: the gate refuses to open for anything else, and an approved call whose re-check comes back DENY doesn't run. There is no path from DENY to a human (AUTHZ-INV-12).

If authorization absorbs approval, you get one opaque step that answers "may this happen?" by some mixture of rules and humans that nobody can separate afterwards. Audit can't say whether a rollback happened because policy permitted it or because someone overrode policy. Changing who approves means changing policy, and changing policy means re-arguing who approves.

Three properties keep the boundary clean:

1. **Approval is bound to one exact call.** It is not attached to the English intention "roll back payment-service". It is bound to a fingerprint of the principal, acting-for, action, resource, environment, incident and arguments, plus the policy version and decision that required it, an

expiry, and a single use. Change any of them and the approval no longer applies: approving `v4.17.2` doesn't approve `v4.16.0`, and an approval given under one policy version doesn't survive a policy change.

- 2. **The approver is authorized too.** Approval needs a human principal holding the role the decision named (`IncidentCommander`), who isn't the requester, within the deadline (`10m`).
- 3. **Authorization runs again at time of use.** The approval authorizes the call. The re-check makes sure the world still agrees.

Everything else about approval is the next article's subject: *when* to require it, how to present the evidence, how to avoid rubber-stamping, what happens when nobody answers, and how humans intervene in a running agent rather than only at a gate.

Auditability and Governance

A decision that isn't recorded may as well not have happened. For every tool call, the platform's decision log should be able to answer the following on its own, without traces, chat history or anyone's memory:

Question	Field	From the POC's rollback
Who acted?	<code>principal</code>	<code>incident-agent-prod</code> (<code>agent.incident-intel@1.3.0</code>)
On whose behalf?	<code>acting_for</code> , <code>delegation_chain</code>	<code>sre-team</code> , delegation <code>dlg-sre-incident-2026q4</code> granted by <code>sre.lead</code> , over a service sre-team owns
What was attempted?	<code>action</code> , <code>resource</code> , <code>environment</code> , <code>arguments</code>	rollback payment-service production → v4.17.2
Which policy was evaluated?	<code>policy_version</code>	<code>authz-2026-09-29.2</code>
Which attributes were used?	<code>attributes</code>	SEV-1, open, normal window, high risk, evidence 0.94 (four observed signals), tier-0
Where did each fact come from?	<code>attribute_sources</code>	severity and state from the incident system, the window from the change calendar, evidence from the evidence evaluator; nothing from the agent
Which rules decided it?	<code>grants</code> , <code>matched_rules</code>	role, delegation, relation; <code>A1-production-rollback</code>
What was the decision?	<code>decision</code> , <code>reason</code> , <code>constraints</code>	ALLOW_WITH_APPROVAL, IncidentCommander, 10m
What was the agent told?	<code>agent_view</code>	<code>APPROVAL_REQUIRED</code> , "Wait for the approval decision"

Question	Field	From the POC's rollback
Was human approval required and given?	<code>approval_id</code> , <code>approved_by</code> , <code>request_fingerprint</code>	<code>apr-819ba4aa10</code> , <code>ic.dev</code> , for fingerprint <code>8fc070d3...</code> (agent's self-approval rejected)
What actually executed?	<code>tool.executed</code> record	rolled back to v4.17.2 at 14:09:02, after re-check
With which credential?	<code>credential</code>	a single-use credential issued for that one call, expiring at 14:10:02

Two fields are easy to skip and expensive to lack.

The policy version. Policy changes. When someone asks in a review why the agent was allowed to do X last Tuesday, "the current policy wouldn't allow that" is not an answer. You need to know what the policy *was*, which means versioning policy like code and stamping every decision with the version that made it.

The attributes. A decision is a function of its inputs. If you log only the output, you can't tell whether a surprising ALLOW came from a bad rule or a bad fact, such as a stale severity or a mislabelled environment. Logging the attributes turns "why did it do that?" into a lookup.

Where the log lives. A hash chain in the agent platform's own storage makes tampering evident; it doesn't make it impossible for whoever runs that storage. Ship the trail as it's written to an independently controlled, append-only store.

This log is the foundation for what comes later in the series. Observability will aggregate it: denial rates per agent, approval latency, which rules fire most. Governance will query it: which agents touched production this quarter, under whose authority. The AI control plane will manage the policy that produces it. None of those are possible if decisions evaporate after they're made.

Production Design Principles

Distilled from the above, as rules you can hold a design review to:

1. **Deny by default.** An action, principal or resource the policy doesn't know is denied. There is no implicit allow anywhere.
2. **Least privilege, by construction.** Agent roles are ceilings sized to the agent's job. A read-only agent has a read-only role, not a shared one.
3. **Narrow scope at four levels.** Tool, action, resource, context. An allowlist of tools is the first level, not the whole answer.
4. **Explicit grants only.** Every grant is a named record with an owner: a role assignment, a relationship, a delegation. If you can't point to the record, it doesn't exist.

5. **Short-lived, conditional authority.** Delegations expire and are tied to conditions (an open incident). Tool credentials live for minutes, as in the last article.
6. **Context rules only restrict.** ABAC can deny, constrain or require approval. It can never grant what no role grants.
7. **Unknown context fails closed.** A missing attribute makes a decision stricter, never looser.
8. **Evidence, never self-assessment.** Any attribute that can loosen a decision is looked up or observed by the platform. The agent doesn't grade its own homework.
9. **One non-bypassable enforcement point.** Every tool call crosses the gateway. The runtime holds no credentials, and network policy leaves no side door, not even for "just this one agent".
10. **Approval gates for dangerous operations, separate from policy.** Decided by policy, performed by a human, bound to one call, and followed by a re-check.
11. **Policy is code.** Versioned, reviewed, tested, deployed independently of agents.
12. **Every decision is explainable and logged, to two audiences.** The agent gets a code, a reason and a next step, but only what it's allowed to learn. The auditor gets everything, in a record shipped to a store the platform can't quietly rewrite.

T2 · AUTHORIZATION & POLICY · THE CHECKLIST

The production checklist.

Twelve rules to hold every agent design review to.

GRANT 01-04	DECIDE 05-08	OPERATE 09-12
01 ☒ Deny by default	05 ☒ Short-lived, conditional authority	09 ☒ One non-bypassable enforcement point
02 ☒ Least-privilege roles, sized to the job	06 ☒ Context rules only restrict	10 ☒ Approval separate from policy
03 ☒ Scope at four levels: tool, action, resource, context	07 ☒ Unknown context fails closed	11 ☒ Policy is code: versioned and tested
04 ☒ Explicit, named grants only	08 ☒ Evidence, never self-assessment	12 ☒ Every decision explained and logged

Hold every agent design review to these twelve.

Production AI Engineering · T2 · Authorization & Policy

Figure 15. The twelve principles on one card, to pin next to your design-review template.

Common Anti-Patterns

Each of these is common, each seems reasonable when it's introduced, and each violates one of the six distinctions.


Anti-pattern	What it looks like	Why it fails	Instead
Admin access "for now"	The agent's credential is cluster-admin, because the first demo needed it	Every prompt-injected instruction becomes an admin action. Nobody ever narrows it later	Role ceilings sized to the job, from day one
Permissions in the prompt	<i>"You must never delete namespaces."</i>	The prompt is input, and input can be overridden: by injection, by a stale runbook line, by a confident model	Put the boundary where the prompt can't reach: the PEP
Tool-level allowlists only	<code>tools: [kubernetes, jira, slack]</code>	Grants every action on every resource in every context	Action, resource and context levels on top
No approval for destructive production actions	The agent can roll back production directly because "it's usually right"	"Usually" is not a control. The one time it's wrong is the incident review	ALLOW_WITH_APPROVAL for high-risk production writes
Untracked delegated authority	The agent acts with a human's token, or with "the platform's" authority	No one can say whose authority was spent. Confused deputies go undetected	Explicit acting-for records, intersected, logged per call
Unlogged policy decisions	The gateway blocks or allows and moves on	Denials can't be tuned, allows can't be justified, incidents can't be reconstructed	Log every decision with version, attributes, rules and reason
Approval merged with authorization	One opaque "can this run?" step mixing rules and humans	Approval becomes a policy bypass. Audit can't tell policy from override	Two gates: policy decides validity, approval decides proceeding
Checks hard-coded in each agent	<code>if env == "prod"</code> in six repositories	Inconsistent, untestable, invisible to anyone but the author	Shared PDP, one PEP, policy as code
Self-graded confidence	Production writes gated on the model's own confidence score	The number moves with phrasing and calibration; the agent grades its own homework	Evidence the platform observed: release correlation, log signature, staging validation, recovery
A bypassable gateway	The runtime also has a kubeconfig "for debugging"	Every control in the gateway becomes advisory	No credentials in the runtime; network policy allows tools only via the gateway
Authorize once per session	The agent is authorized at start and then trusted for the run	Multi-step agents change what they're doing; context changes under them	Authorize every call; re-check after any wait

The Final Architecture View

Put together, the platform now has two control layers in front of every tool, with a third, dashed in the figure, to come:

Who · whether · when · remember.

Identity says who. Policy says whether. Approval says when a human must agree. Audit remembers all three.



Identity says who. Policy says whether. Approval says when. Audit remembers all three.

Figure 16. Identity says who. Policy says whether. Approval says when a human must agree. Audit remembers all three.

Reading left to right, following the production rollback:

- Ingress and identity (T1).** The Datadog event is authenticated. The execution gets an identity: `incident-agent-prod`, acting for `sre-team`.
- Reasoning (F2, F3).** The agent proposes `rollbackDeployment(payment-service, production, v4.17.2)`. It holds no credentials, and the network gives it no route to Kubernetes except through the gateway.
- Enforcement (F1 gateway, now a PEP).** The gateway builds the authorization request and asks for a decision.
- Decision (T2).** The policy engine combines role, relationship, context and platform-observed evidence, and returns `ALLOW_WITH_APPROVAL`.
- Approval (T3, next).** The incident commander approves this exact call. The agent re-submits it, the fingerprint matches, and the gateway authorizes again.
- Execution and audit.** The token broker mints a short-lived, audience-bound credential for this call (in the POC, a single-use credential that expires a minute after issue). Kubernetes rolls back. Every step is in the chain.

What Policy Cannot Decide

Look again at the one decision in this article that isn't a plain yes or no:

JSON

```
{
  "decision": "ALLOW_WITH_APPROVAL",
  "code": "APPROVAL_REQUIRED",
  "reason": "Production rollback requires SRE approval",
  "next": "Wait for the approval decision",
  "constraints": { "approverRole": "IncidentCommander", "expiresIn": "10m" }
}
```

Policy did everything it can do here. It established that the action is valid: the right kind of agent, spending authority that really exists, on a resource its principal owns, in a context where rollback is reasonable. And it concluded that valid isn't enough, and that **a human must still decide whether this happens now**.

That's where authorization stops. It can say *that* a human is needed and *which* human. It can't say how that human should decide, what evidence they need in front of them, how long the agent should wait, what happens if the approver says "not yet, try restarting first", or how a human steps into a running agent that's heading somewhere unexpected without waiting to be asked.




Figure 17. Identity told us who is acting. Policy told us what they may do. Next: when a human must decide anyway. Those are the questions of the next control layer.

Policy may say an action is valid. Execution may still require explicit human approval. The next article is Human-in-the-Loop.

Closing

Back to 14:09, and the question Agent Identity left open.

Should incident-agent-prod be allowed to roll back payment-service in production, at 14:09, during a SEV-1, with a change freeze starting at 15:00? Who decides?

We can now answer it precisely. **Yes, conditionally.** The role `incident-responder` permits rollbacks. sre-team delegated rollback authority for incidents on services it owns, and it owns payment-service. The incident is open and SEV-1. The diagnosis evidence clears the 0.80 floor at 0.94, and the platform observed every signal behind it: a release thirteen minutes before the alert, logs naming that release, a validated staging rollback and a recovered staging environment. The freeze hasn't started, and even if it had, a SEV-1 would still qualify. So policy `authz-2026-09-29.2` returned `ALLOW_WITH_APPROVAL` under rule `A1-production-rollback`. The incident commander approved that exact call, fingerprint `8fc070d3...`, at 14:08:50. The agent re-submitted it, the gateway checked again at 14:09:02, and the rollback executed.

Who decides? Not the agent: it proposes, and its opinion of itself never enters the decision. Not the tool: it only sees a credential. Not the prompt: it's input. A policy engine decides whether the action is valid. A named human decides whether it proceeds. The gateway enforces both. And the log remembers every step.

Agent systems need policy decisions that are contextual, auditable and enforceable at runtime. Identity tells us who is acting. Authorization tells us what they may do. Approval tells us when a human must still say yes.

References

- OWASP, Top 10 for LLM Applications 2025: LLM06 Excessive Agency
- NIST, SP 800-162: Guide to Attribute Based Access Control (ABAC)
- NIST, SP 800-207: Zero Trust Architecture: policy decision and enforcement points
- OASIS, XACML 3.0: PEP / PDP / PIP / PAP, obligations and advice
- Pang et al., Zanzibar: Google's Consistent, Global Authorization System, USENIX ATC 2019: relationship-based access control

- Cedar policy language: default deny, forbid-overrides-permit, policy annotations
 - Open Policy Agent and OpenFGA
 - IETF, RFC 8693: OAuth 2.0 Token Exchange: delegation vs impersonation
 - Model Context Protocol, Authorization and Security best practices
 - Kubernetes, Using RBAC Authorization
 - OpenFGA, Modeling agents and acting on behalf of: relationship-based authorization for agents
 - Norm Hardy, *The Confused Deputy (or why capabilities might have been invented)*, ACM SIGOPS Operating Systems Review (1988)
-

Series. F1 · MCP Tool Sprawl · F2 · Layered Architecture · F3 · Headless AI · T1 · Agent Identity · **T2 · Authorization & Policy** (this article) · Next: T3 · Human-in-the-Loop · then AI Control Plane · Observability & Governance · Production Agent Platform.

Every decision, identifier, timestamp and measured value in this article comes from

`authz_poc/runs/2026-09-29-recorded/` (`decisions.jsonl` , `facts.json` , `timeline.json` , `sweep.json` , `receipt.json` , `expectations.json` , `invariants.json` , `tests.json` , `verification.json`). To reproduce: `cd authz_poc && python3 -m authz.run && python3 -m unittest discover -s tests -v && python3 -m authz.verify` . The claim-by-claim map is `docs/TRACEABILITY.md` ; the Lab Console is `results/lab-console.html` (`python3 tools/build_lab_console.py`).